

Wrangling_data

March 27, 2022

1 Predict Heart Disease Status Based on Quantifiable Variables

2 Introduction:

Cardiovascular diseases (CVDs) is a class of disease that involves the heart or blood vessels. the number one cause of death globally, taking an estimated 17.9 million lives each year, which accounts for 31% of all deaths worldwide. Four out of five CVD deaths are due to heart attacks and strokes, and one-third of these deaths occur prematurely in people under 70 years of age. Heart failure is a common event caused by CVDs.

People with cardiovascular disease or who are at high cardiovascular risk (due to the presence of one or more risk factors such as hypertension, diabetes, hyperlipidemia or already established disease) need early detection and management wherein a machine learning model can be of great help.

Our question is can we determine heart disease status based on quantifiable variables. This dataset we will be using is a tabular data set with comma-separated variables. It has 12 variables, 11 predictors variables and one target variable which is heart disease.

2.0.1 Attribute Information

1. Age: years
2. Sex: (0 = MALE, 1 = FEMALE)
3. ChestPainType: (ATA = 1, NAP = 2, ASY =3, TA =4)
4. RestingBP: resting blood pressure (mm HG)
5. Cholesterol: (mm/dl)
6. Fasting Blood Sugar: fasting blood sugar (1: if fastingBS >120 mg/dlm 0: otherwise)
7. RestingECG: Resting Electrocardiogram Results
8. MaxHR: maximum heart rate achieved (Numeric value between 60 and 202)
9. ExerciseAngina: exercise-induced angina (Y: Yes, N: No)
10. Oldpeak: (Numeric value measured in depression)
11. ST_Slope: the slope of the peak exercise ST segment (Up: upsloping, Flat: flat, Down: downsloping)
12. HeartDisease:(1: heart disease, 0: Normal)

```
[441]: set.seed(8)
library(repr)
library(tidyverse)
library(tidymodels)
library(dplyr)
```

```
library(RColorBrewer)
```

2.0.2 Reading Files: Wrangling and Cleaning Data Set

```
[442]: set.seed(8)
heart_data <- read_csv("heart.csv") %>%
  mutate(HeartDisease = as_factor(HeartDisease))

head(heart_data)
```

Parsed with column specification:

```
cols(
  Age = col_double(),
  Sex = col_character(),
  ChestPainType = col_character(),
  RestingBP = col_double(),
  Cholesterol = col_double(),
  FastingBS = col_double(),
  RestingECG = col_character(),
  MaxHR = col_double(),
  ExerciseAngina = col_character(),
  Oldpeak = col_double(),
  ST_Slope = col_character(),
  HeartDisease = col_double()
)
```

A tibble: 6 × 12

	Age <dbl>	Sex <chr>	ChestPainType <chr>	RestingBP <dbl>	Cholesterol <dbl>	FastingBS <dbl>	RestingECG <chr>	MaxHR <dbl>
	40	M	ATA	140	289	0	Normal	172
	49	F	NAP	160	180	0	Normal	156
	37	M	ATA	130	283	0	ST	98
	48	F	ASY	138	214	0	Normal	108
	54	M	NAP	150	195	0	Normal	122
	39	M	NAP	120	339	0	Normal	170

Split data to train and test¶

```
[443]: set.seed(8)
heart_split <- initial_split(heart_data, prop = 0.75, strata = HeartDisease)
heart_train <- training(heart_split)
heart_test <- testing(heart_split)
```

2.0.3 Summarize dataset

```
[444]: set.seed(8)
num_obs <- nrow(heart_train)
heart_sum <- heart_train %>%
  glimpse() %>%
  group_by(HeartDisease) %>%
  summarize(count = n(), percentage = n()/ num_obs* 100)

heart_sum

checking_for_na <- sum(is.na(heart_train))
checking_for_na

summary(heart_train)
```

```
Rows: 689
Columns: 12
$ Age          <dbl> 40, 49, 37, 48, 39, 45, 37, 58, 39,
49, 42, 54, 38, 60...
$ Sex          <chr> "M", "F", "M", "F", "M", "F", "M",
"M", "M", "M", "F",...
$ ChestPainType <chr> "ATA", "NAP", "ATA", "ASY", "NAP",
"ATA", "ASY", "ATA"...
$ RestingBP    <dbl> 140, 160, 130, 138, 120, 130, 140,
136, 120, 140, 115,...
$ Cholesterol  <dbl> 289, 180, 283, 214, 339, 237, 207,
164, 204, 234, 211,...
$ FastingBS    <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, ...
$ RestingECG   <chr> "Normal", "Normal", "ST", "Normal",
"Normal", "Normal"...
$ MaxHR        <dbl> 172, 156, 98, 108, 170, 170, 130, 99,
145, 140, 137, 1...
$ ExerciseAngina <chr> "N", "N", "N", "Y", "N", "N", "Y",
"Y", "N", "Y", "N",...
$ Oldpeak      <dbl> 0.0, 1.0, 0.0, 1.5, 0.0, 0.0, 1.5,
2.0, 0.0, 1.0, 0.0,...
$ ST_Slope     <chr> "Up", "Flat", "Up", "Flat", "Up",
"Up", "Flat", "Flat"...
$ HeartDisease <fct> 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1,
1, 1, 0, 0, 1, ...

`summarise()` ungrouping output (override with `.groups` argument)
```

	HeartDisease <fct>	count <int>	percentage <dbl>
A tibble: 2 × 3	0	308	44.70247
	1	381	55.29753

0

Age	Sex	ChestPainType	RestingBP
Min. :28.00	Length:689	Length:689	Min. : 0.0
1st Qu.:47.00	Class :character	Class :character	1st Qu.:120.0
Median :54.00	Mode :character	Mode :character	Median :130.0
Mean :53.53			Mean :132.8
3rd Qu.:60.00			3rd Qu.:140.0
Max. :77.00			Max. :200.0
Cholesterol	FastingBS	RestingECG	MaxHR
Min. : 0.0	Min. :0.0000	Length:689	Min. : 63.0
1st Qu.:173.0	1st Qu.:0.0000	Class :character	1st Qu.:120.0
Median :222.0	Median :0.0000	Mode :character	Median :138.0
Mean :198.5	Mean :0.2293		Mean :136.8
3rd Qu.:267.0	3rd Qu.:0.0000		3rd Qu.:156.0
Max. :564.0	Max. :1.0000		Max. :202.0
ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
Length:689	Min. :-2.6000	Length:689	0:308
Class :character	1st Qu.: 0.0000	Class :character	1:381
Mode :character	Median : 0.5000	Mode :character	
	Mean : 0.8856		
	3rd Qu.: 1.5000		
	Max. : 5.6000		

Observations from Summary

1. Resting blood pressure and Cholesterol have zero as a minimum which is unusual.
2. There may be outliers/Missings in Cholesterol and Resting BP being presented as zero.
3. Number of rows 689 and number of columns 12.
4. Percentage of people with heart disease: 44.70 %
5. Percentage of people without heart disease: 55.30%

Fixing zeros in and Resting BP and Cholesterol

```
[445]: set.seed(8)
heart_train <- heart_train%>%
  filter(RestingBP != 0, Cholesterol != 0)
count(heart_train)
summary(heart_train)
head(heart_train)
```

A tibble: 1 × 1

n
562

Age	Sex	ChestPainType	RestingBP
Min. :28.00	Length:562	Length:562	Min. : 92.0
1st Qu.:46.00	Class :character	Class :character	1st Qu.:120.0
Median :54.00	Mode :character	Mode :character	Median :130.0
Mean :52.94			Mean :133.7

3rd Qu.:60.00		3rd Qu.:140.0
Max. :77.00		Max. :200.0
Cholesterol	FastingBS	RestingECG
Min. : 85.0	Min. :0.0000	Length:562
1st Qu.:206.0	1st Qu.:0.0000	Class :character
Median :236.0	Median :0.0000	Mode :character
Mean :243.3	Mean :0.1655	
3rd Qu.:274.8	3rd Qu.:0.0000	
Max. :564.0	Max. :1.0000	
ExerciseAngina	Oldpeak	ST_Slope
Length:562	Min. :0.0000	Length:562
Class :character	1st Qu.:0.0000	Class :character
Mode :character	Median :0.5000	Mode :character
	Mean :0.9192	
	3rd Qu.:1.6000	
	Max. :5.6000	

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	Maxi
	<dbl>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<chr>	<dbl>
A tibble: 6 × 12	40	M	ATA	140	289	0	Normal	172
	49	F	NAP	160	180	0	Normal	156
	37	M	ATA	130	283	0	ST	98
	48	F	ASY	138	214	0	Normal	108
	39	M	NAP	120	339	0	Normal	170
	45	F	ATA	130	237	0	Normal	170

2.0.4 Visualisations of data

We will visualise each predictor variable to see how much it factors into a person having heart disease or not. Accuracy of predictor variables isn't directly proportional to the number of variables to use. In our case we have 11 possible predictors and it would be impossible to take them all into account when making a classification. So we analyse each of the eleven variables separately and see its correlation with heart disease by making eleven different graphs.

0: No heart Disease

1: Heart Disease

1. Heart Disease with Age:

```
[446]: set.seed(8)
HeartDisease_Age_plot <- heart_train %>%
  ggplot(aes(x = Age, fill = Age)) +
  geom_bar() +
  facet_grid(~HeartDisease) +
  labs(title = "Heart Disease with Age", x = "Age_
  ↪(Years) ", y = "Number of Patients with heart Disease")
options(repr.plot.width = 11, repr.plot.height = 8)

HeartDisease_Age_plot
```




Figure 1: The graph above explores the relationship between heart disease and age. By looking at the graphs we can see that people between the ages of 55-65 seem to have the most heart diseases. So this seems like a good predictor variable to use.

2. Heart Disease with Sex:

```
[447]: set.seed(8)
HeartDisease_Sex_plot <- heart_train %>%
  ggplot(aes(x = Sex, fill = Sex)) +
  geom_bar() +
  facet_grid(~HeartDisease) +
  geom_text(aes(label = ..count..), stat = "count",
  ↪ vjust = 2, colour = "black") +
  labs(title = "Heart Disease with Sex", x = "Sex", y =
  ↪ "Count") +
  scale_color_manual (labels = c("Female", "Male"))
options(repr.plot.width = 8, repr.plot.height = 8)
HeartDisease_Sex_plot
```




FIGURE 2: There seems to be a strong correlation between sex and heart diseases. As we can see from the graph above it makes sense to use it as one of our predictor variables.

3. Heart Disease vs. Chest Pain Type: (TA: Typical Angina, ATA: Atypical Angina, NAP: Non-Anginal Pain, ASY: Asymptomatic)

```
[448]: set.seed(8)
HeartDisease_ChestPainType_plot <- heart_train %>%
  ggplot(aes(x = ChestPainType, fill = 
    ↪ChestPainType)) +
  geom_bar() +
  facet_grid(~HeartDisease) +
```

```

    geom_text(aes(label = ..count..), stat = "count", vjust = 2, colour = "black") +
    labs(title = "Heart Disease with Chest Pain Type", x = "Chest Pain Type", y = "Count",
         fill = "Chest Pain Type")
HeartDisease_ChestPainType_plot

```

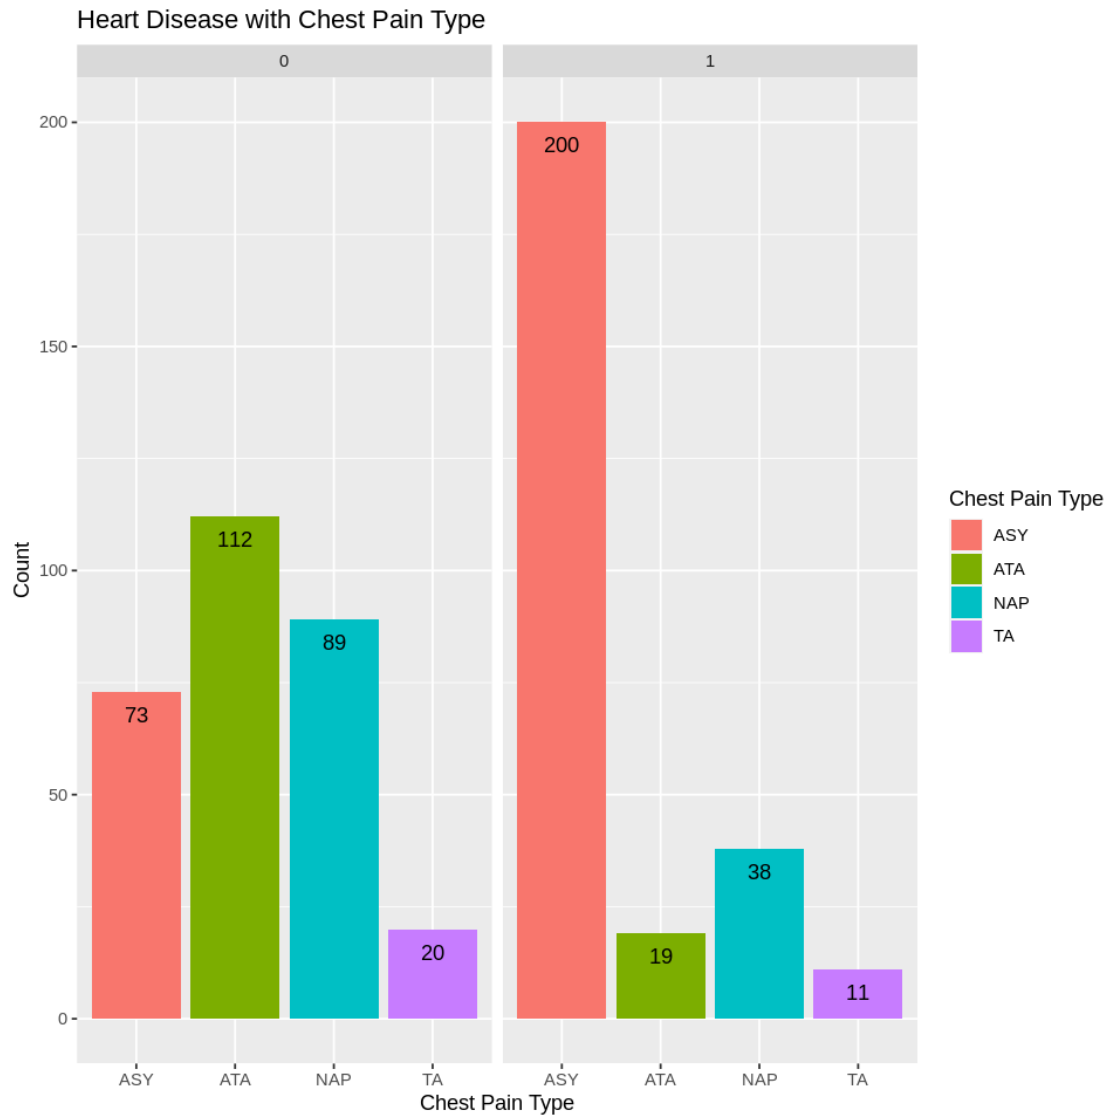


Figure 3: The above graph shows the relationship between the type of chest pain and heart disease. People with ASY chest pain seem to be more likely to have heart disease.

4. Heart Disease vs. Resting Blood Pressure (mm HG):


```
[449]: set.seed(8)
HeartDisease_RestingBP_plot <- heart_train %>%
  ggplot(aes(x = RestingBP, fill = RestingBP)) +
  geom_bar() +
  facet_grid(~HeartDisease) +
  labs(title = "Heart Disease with Resting Blood
  ↪Pressure", x = "Resting Blood Pressure", y = "Count")
options(repr.plot.width = 12, repr.plot.height = 8)

HeartDisease_RestingBP_plot
```

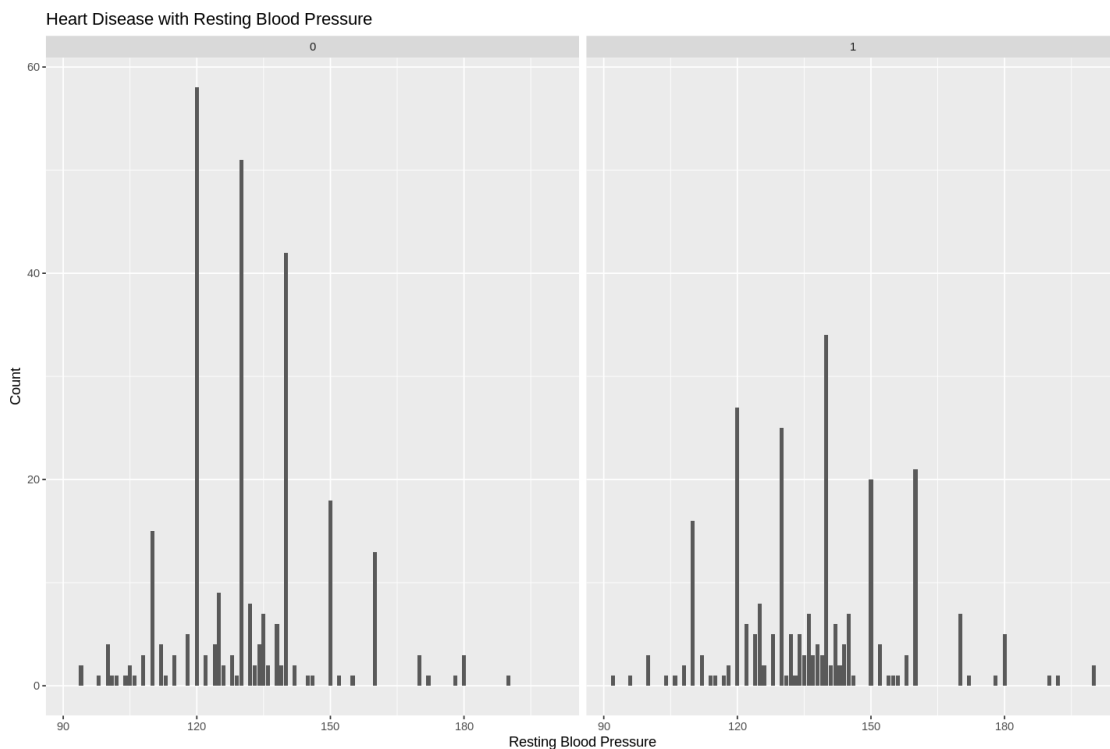


Figure 4: There doesn't seem to be any strong correlation between the resting blood pressure of a person and heart disease.

5. Heart Disease vs. Cholesterol (mm/dl):

```
[450]: set.seed(8)
HeartDisease_Cholesterol_plot <- heart_train %>%
  ggplot(aes(x = Cholesterol, fill =
  ↪Cholesterol)) +
  facet_grid(~HeartDisease) +
  geom_bar() +
  labs(title = "Heart Disease with Serum
  ↪Cholestero", x = "Serum Cholestero", y = "Count")
```

```
options(repr.plot.width = 8, repr.plot.height = 8)
```

```
HeartDisease_Cholesterol_plot
```

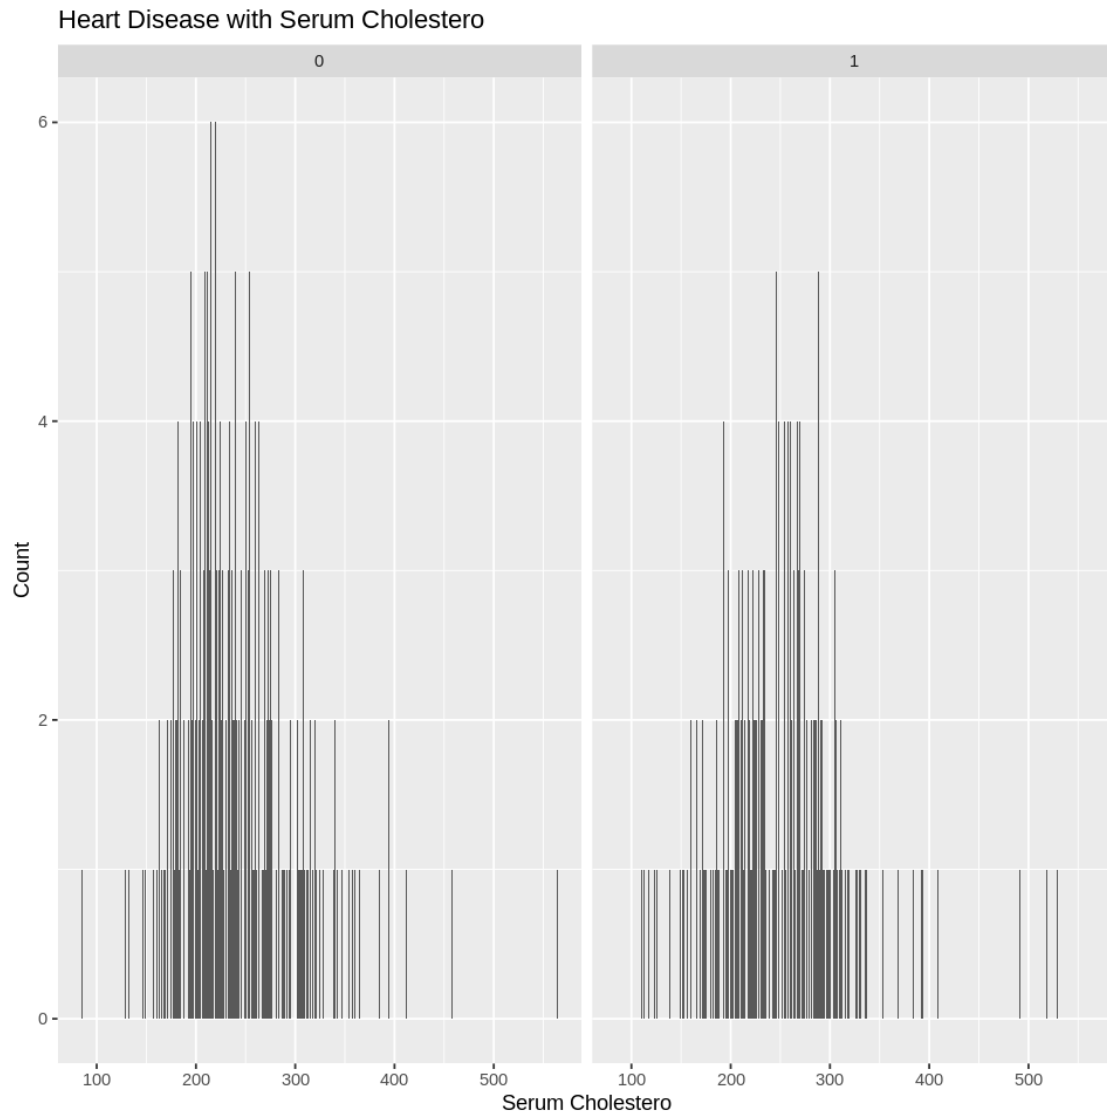


Figure 5: There doesn't seem to be any kind of relationship between cholesterol and heart disease, so it doesn't make sense to use it as a predictor variable.

6. Heart Disease vs. Fasting Blood Sugar:

```
[451]: set.seed(8)
HeartDisease_FastingBS_plot <- heart_train %>%
  ggplot(aes(x = FastingBS, fill = as.
    ↪character(FastingBS))) +
```

```

geom_bar() +
facet_grid(~HeartDisease) +
geom_text(aes(label = ..count..), stat = "count", vjust = 2, colour = "white") +
labs(title = "Heart Disease with Fasting Blood Sugar",
x = "Fasting Blood Sugar", y = "Count",
fill = "Fasting Blood Sugar")
HeartDisease_FastingBS_plot

```

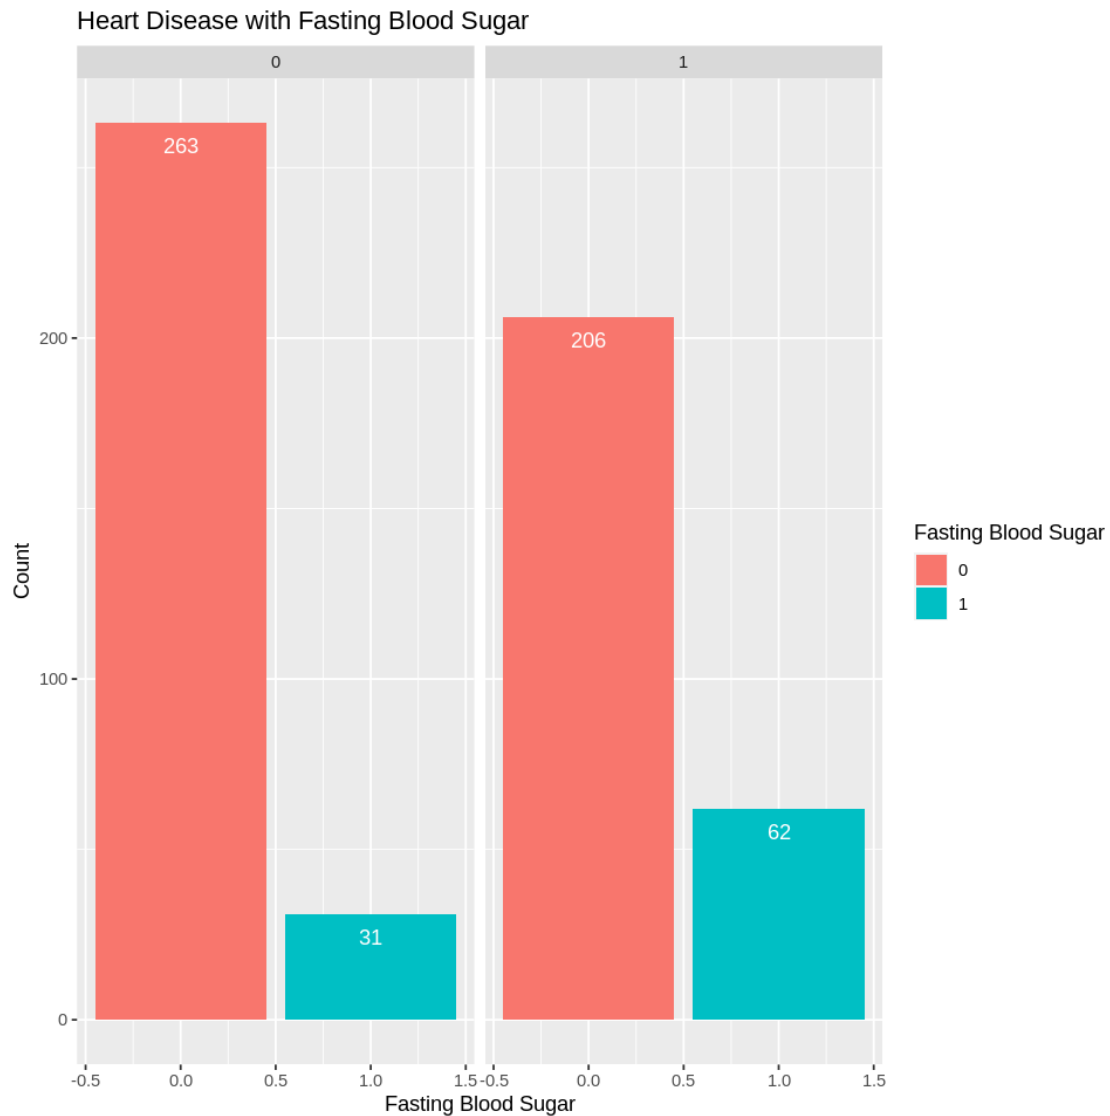


Figure 6: There doesn't seem to be any kind of relationship between fasting blood sugar and heart disease, so it doesn't make sense to use it as a predictor variable.

7. Heart Disease vs. Resting Electrocardiogram Results: (Normal: Normal, ST: having ST-T wave abnormality, LVH: showing probable or definite left ventricular hypertrophy by Estes' criteria)

```
[452]: set.seed(8)
HeartDisease_RestingECG_plot <- heart_train %>%
  ggplot(aes(x = RestingECG, fill = RestingECG)) +
  geom_bar() +
  facet_grid(~HeartDisease) +
  geom_text(aes(label = ..count..), stat = "count",
    vjust = 2, colour = "black") +
  labs(title = "Heart Disease with Resting Electrocardiogram Results",
    x = "Resting Electrocardiogram Results", y = "Count",
    fill = "Resting Electrocardiogram Results")
HeartDisease_RestingECG_plot
```




Figure 7: There seems to be a weak positive relationship between the Resting ECG and whether a person has heart disease or not. People with a normal resting ECG are more likely to face cardiovascular problems. However there's not a strong enough relationship to make it a predictor variable.

8. Heart Disease vs. Maximum Heart Rate: Maximum Heart Rate: Numeric value between 60 and 202

```
[453]: set.seed(8)
HeartDisease_RestingBP_plot <- heart_train %>%
  ggplot(aes(x = MaxHR, fill = MaxHR)) +
  facet_grid(~HeartDisease) +
  geom_bar() +
```

```

labs(title = "Heart Disease with Maximum Heart_Rate", x = "Maximum Heart Rate", y = "Count")
options(repr.plot.width = 8, repr.plot.height = 8)

```

HeartDisease_RestingBP_plot

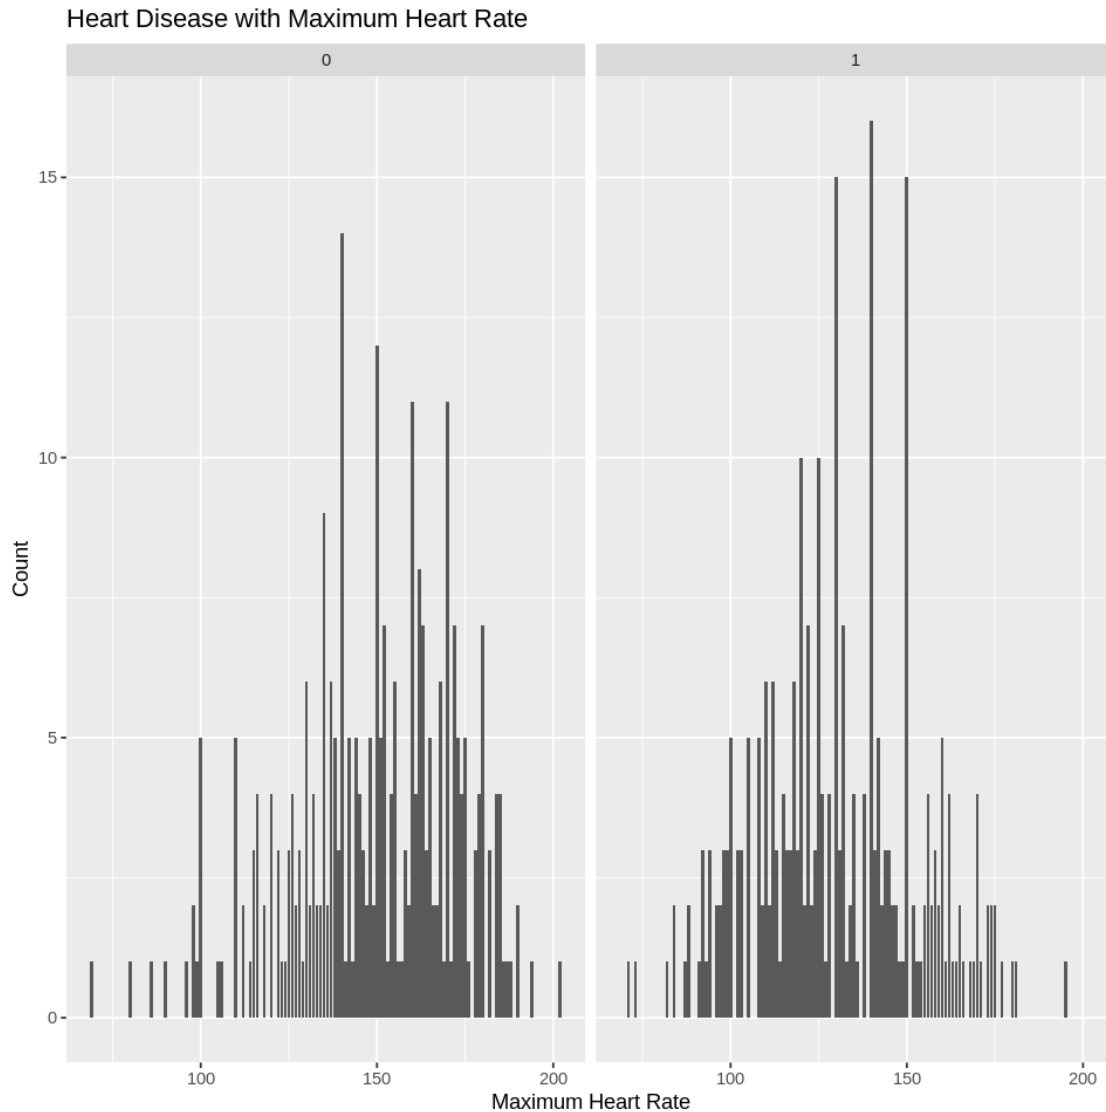


Figure 8: People with a maximum heart rate in the range of 100-150 are more likely to have heart disease. However there's an overlap between people who have heart disease and those who don't have heart in the same range. It increases and decreases rapidly instead of going in one direction. So it doesn't make sense to use it as a predictor variable.

9. Heart Disease vs. Exercise-Induced Angina: Y: The person has Exercise-Induced Angina
N: The person doesn't have Exercise-Induced Angina

```
[454]: set.seed(8)
HeartDisease_ExerciseAngina_plot <- heart_train %>%
  ggplot(aes(x = ExerciseAngina, fill =
    ↪ExerciseAngina)) +
    geom_bar() +
    facet_grid(~HeartDisease) +
    geom_text(aes(label = ..count..), stat =
    ↪"count", vjust = 2, colour = "black") +
    labs(title = "Heart Disease with
    ↪Exercise-Induced Angina",
    x = "Exercise-Induced Angina", y =
    ↪"Count",
    fill = "Exercise-Induced Angina")
HeartDisease_ExerciseAngina_plot
```




Figure 9: There seems to be a strong relationship between exercise-Induced Angina and Heart Disease. People with no exercise angina are more likely to have heart disease.

10. Heart Disease vs. Old peak: Old peak is numeric value measured in depression.

```
[455]: set.seed(8)
HeartDisease_RestingBP_plot <- heart_train %>%
  ggplot(aes(x = Oldpeak, fill = Oldpeak)) +
  facet_grid(~HeartDisease) +
  geom_bar() +
  labs(title = "Heart Disease with Old peak", x = "Old peak", y = "Count")
options(repr.plot.width = 8, repr.plot.height = 8)

HeartDisease_RestingBP_plot
```





Figure 10: There's a weak relationship between old peka and heart disease, but it's not strong enough to make it a predictor variable.

11. Heart Disease vs. ST_Slope: (ST_Slope: the slope of the peak exercise ST segment Up: upsloping, Flat: flat, Down: downsloping)

```
[456]: set.seed(8)
HeartDisease_ST_Slope_plot <- heart_train %>%
  ggplot(aes(x = ST_Slope, fill = HeartDisease)) +
  geom_bar() +
  facet_grid(~HeartDisease) +
  geom_text(aes(label = ..count..), stat = "count",
  ↪ vjust = 2, colour = "black") +
```

```
labs(title = "Heart Disease with ST Slope", x = "ST Slope", y = "Count", fill = "Heart Disease")
```

HeartDisease_ST_Slope_plot




Figure 11: There's a relationship between ST_slope and heart disease. If the St_slope is up then the person is more likely to have a heart disease, whereas if the St_slope is flat then the person is more likely to.

2.0.5 Summary of Exploratory data analysis:

From the above 11 graphs we can conclude that we will only have five predictor variables: age, sex, ExerciseAngina, ST_Slope and chest pain type. Generally accuracy of a predictor model is not

directly proportional to the number of variables that we use. So it wouldn't make sense to use all 11 variables as our predictors, which is why we narrowed down our choices.

3 Data Analysis:

Since we have our five predictor variables and one target variable, we will now make a prediction model to classify future patients and determine whether they have heart disease or not.

Convert Characters to Numerics We convert the columns of sex, chest Pain type, ST_Slope and Exercise Angina to numerical values as we can't standardize data with characters. The other predictor variable age is already stored as numeric value so no need to convert it.

```
[457]: set.seed(8)
heart_train <- heart_train %>%
  mutate(Sex = as_factor(Sex)) %>%
  mutate(Sex = as.numeric(Sex)) %>%
  mutate(ChestPainType = as_factor(ChestPainType)) %>%
  mutate(ChestPainType = as.numeric(ChestPainType)) %>%
  mutate(ExerciseAngina = as_factor(ExerciseAngina)) %>%
  mutate(ExerciseAngina = as.numeric(ExerciseAngina)) %>%
  mutate(ST_Slope = as_factor(ST_Slope)) %>%
  mutate(ST_Slope = as.numeric(ST_Slope))
head(heart_train)

heart_test <- heart_test %>%
  mutate(Sex = as_factor(Sex)) %>%
  mutate(Sex = as.numeric(Sex)) %>%
  mutate(ChestPainType = as_factor(ChestPainType)) %>%
  mutate(ChestPainType = as.numeric(ChestPainType)) %>%
  mutate(ExerciseAngina = as_factor(ExerciseAngina)) %>%
  mutate(ExerciseAngina = as.numeric(ExerciseAngina)) %>%
  mutate(ST_Slope = as_factor(ST_Slope)) %>%
  mutate(ST_Slope = as.numeric(ST_Slope))
head(heart_test)
```

	Age <dbl>	Sex <dbl>	ChestPainType <dbl>	RestingBP <dbl>	Cholesterol <dbl>	FastingBS <dbl>	RestingECG <chr>	Maxi <dbl>
A tibble: 6 × 12	40	1	1	140	289	0	Normal	172
	49	2	2	160	180	0	Normal	156
	37	1	1	130	283	0	ST	98
	48	2	3	138	214	0	Normal	108
	39	1	2	120	339	0	Normal	170
	45	2	1	130	237	0	Normal	170

	Age <dbl>	Sex <dbl>	ChestPainType <dbl>	RestingBP <dbl>	Cholesterol <dbl>	FastingBS <dbl>	RestingECG <chr>	MaxHR <dbl>
A tibble: 6 × 12	54	1	1	150	195	0	Normal	122
	54	1	2	110	208	0	Normal	142
	48	2	2	120	284	0	Normal	120
	37	2	1	130	211	0	Normal	142
	43	2	2	120	201	0	Normal	165
	43	2	3	100	223	0	Normal	142

Finding the Best K

Use Best K to Predict

```
[458]: set.seed(8)
knn_tune <- nearest_neighbor(weight_func = "rectangular", neighbors = tune())%>%
set_engine("kkn") %>%
set_mode("classification")

heart_recipe <- recipe(HeartDisease ~ Age + Sex + ChestPainType + MaxHR, data_
  ↪=heart_train) %>%
step_scale(all_predictors()) %>%
step_center(all_predictors())

preprocessed_data <- heart_recipe %>%
prep() %>%
bake(heart_train)

heart_vfold <- vfold_cv(heart_train, v = 5, strata = HeartDisease)

heart_workflow <- workflow() %>%
add_recipe(heart_recipe) %>%
add_model(knn_tune)

gridvals <- tibble(neighbors = seq(from = 1, to = 20))

heart_results <- heart_workflow %>%
tune_grid(resamples = heart_vfold, grid = gridvals) %>%
collect_metrics()

heart_min <- heart_results %>%
filter(.metric == "accuracy") %>%
arrange(mean) %>%
head(n = 1)

accuracy_versus_k <- ggplot(heart_results, aes(x = neighbors, y = mean))+
geom_point() +
geom_line() +
labs(x = "Neighbors", y = "Accuracy Estimate")
```

```

k_min <- heart_min %>%
pull(neighbors)

knn_best <- nearest_neighbor(weight_func = "rectangular", neighbors = k_min) %>%
set_engine("kkn") %>%
set_mode("classification")

heart_best_fit <- workflow() %>%
add_recipe(heart_recipe) %>%
add_model(knn_best) %>%
fit(data = heart_train)

heart_test <- heart_test%>%
mutate(Sex = as_factor(Sex)) %>%
mutate(Sex = as.numeric(Sex)) %>%
mutate(ChestPainType = as_factor(ChestPainType)) %>%
mutate(ChestPainType = as.numeric(ChestPainType))

heart_test_predictions <- predict(heart_best_fit, heart_test) %>%
bind_cols(heart_test)

heart_summary <- heart_test_predictions %>%
metrics(truth = HeartDisease, estimate = .pred_class)

heart_summary

```

A tibble: 2 × 3

	.metric <chr>	.estimator <chr>	.estimate <dbl>
	accuracy	binary	0.6462882
	kap	binary	0.3022232

```

[460]: set.seed(8)
knn_tune <- nearest_neighbor(weight_func = "rectangular", neighbors = tune())
  ↳ %>%
set_engine("kkn") %>%
set_mode("classification")

heart_recipe <- recipe(HeartDisease ~ ST_Slope + Sex + ChestPainType +
  ↳ MaxHR, data = heart_train) %>%
step_scale(all_predictors()) %>%
step_center(all_predictors())

heart_vfold <- vfold_cv(heart_train, v = 5, strata = HeartDisease)
heart_workflow <- workflow() %>%
add_recipe(heart_recipe) %>%

```

```

add_model(knn_tune)

gridvals <- tibble(neighbors = seq(from = 1, to = 20))

heart_results <- heart_workflow %>%
  tune_grid(resamples = heart_vfold, grid = gridvals) %>%
  collect_metrics()

heart_min <- heart_results %>%
  filter(.metric == "accuracy") %>%
  arrange(mean) %>%
  head(n = 1)

accuracy_versus_k <- ggplot(heart_results, aes(x = neighbors, y = mean))+
  geom_point() +
  geom_line() +
  labs(x = "Neighbors", y = "Accuracy Estimate")

k_min <- heart_min %>%
  pull(neighbors)

knn_best <- nearest_neighbor(weight_func = "rectangular", neighbors = k_min) %>%
  set_engine("kknn") %>%
  set_mode("classification")

heart_best_fit <- workflow() %>%
  add_recipe(heart_recipe) %>%
  add_model(knn_best) %>%
  fit(data = heart_train)

heart_test_predictions <- predict(heart_best_fit, heart_test) %>%
  bind_cols(heart_test)

heart_summary <- heart_test_predictions %>%
  metrics(truth = HeartDisease, estimate = .pred_class)
heart_summary

```

A tibble: 2 × 3

	.metric <chr>	.estimator <chr>	.estimate <dbl>
	accuracy	binary	0.7467249
	kap	binary	0.4989437

```

[461]: set.seed(8)
knn_tune <- nearest_neighbor(weight_func = "rectangular", neighbors = tune())
  ↪ %>%
set_engine("kknn") %>%
set_mode("classification")

```

```

heart_recipe <- recipe(HeartDisease ~ Age + Cholesterol + Oldpeak + MaxHR +
  ↪FastingBS, data = heart_train) %>%
step_scale(all_predictors()) %>%
step_center(all_predictors())

heart_vfold <- vfold_cv(heart_train, v = 5, strata = HeartDisease)

heart_workflow <- workflow() %>%
add_recipe(heart_recipe) %>%
add_model(knn_tune)

gridvals <- tibble(neighbors = seq(from = 1, to = 20))

heart_results <- heart_workflow %>%
tune_grid(resamples = heart_vfold, grid = gridvals) %>%
collect_metrics()

heart_min <- heart_results %>%
filter(.metric == "accuracy") %>%
arrange(mean) %>%
head(n = 1)

accuracy_versus_k <- ggplot(heart_results, aes(x = neighbors, y = mean))+
geom_point() +
geom_line() +
labs(x = "Neighbors", y = "Accuracy Estimate")

k_min <- heart_min %>%
pull(neighbors)

knn_best <- nearest_neighbor(weight_func = "rectangular", neighbors = k_min) %>%
set_engine("kkn") %>%
set_mode("classification")

heart_best_fit <- workflow() %>%
add_recipe(heart_recipe) %>%
add_model(knn_best) %>%
fit(data = heart_train)

heart_test <- heart_test %>%
mutate(ST_Slope = as_factor(ST_Slope)) %>%
mutate(ST_Slope = as.numeric(ST_Slope))

heart_test_predictions <- predict(heart_best_fit, heart_test) %>%
bind_cols(heart_test)

```

```
heart_summary <- heart_test_predictions %>%
metrics(truth = HeartDisease, estimate = .pred_class)
heart_summary
```

A tibble: 2 × 3

	.metric <chr>	.estimator <chr>	.estimate <dbl>
	accuracy	binary	0.6724891
	kap	binary	0.3490126

```
[462]: set.seed(8)
knn_tune <- nearest_neighbor(weight_func = "rectangular", neighbors = tune())
  ↳ %>%
set_engine("kkn") %>%
set_mode("classification")

heart_recipe <- recipe(HeartDisease ~ ExerciseAngina + ST_Slope + Sex +
  ↳ ChestPainType, data = heart_train) %>%
step_scale(all_predictors()) %>%
step_center(all_predictors())

heart_vfold <- vfold_cv(heart_train, v = 5, strata = HeartDisease)
heart_workflow <- workflow() %>%
add_recipe(heart_recipe) %>%
add_model(knn_tune)

gridvals <- tibble(neighbors = seq(from = 1, to = 20))

heart_results <- heart_workflow %>%
tune_grid(resamples = heart_vfold, grid = gridvals) %>%
collect_metrics()

heart_min <- heart_results %>%
filter(.metric == "accuracy") %>%
arrange(mean) %>%
head(n = 1)

accuracy_versus_k <- ggplot(heart_results, aes(x = neighbors, y = mean))+
geom_point() +
geom_line() +
labs(x = "Neighbors", y = "Accuracy Estimate")

k_min <- heart_min %>%
pull(neighbors)

knn_best <- nearest_neighbor(weight_func = "rectangular", neighbors = k_min) %>%
set_engine("kkn") %>%
set_mode("classification")
```



```

heart_best_fit <- workflow() %>%
add_recipe(heart_recipe) %>%
add_model(knn_best) %>%
fit(data = heart_train)

heart_test_predictions <- predict(heart_best_fit, heart_test) %>%
bind_cols(heart_test)

heart_summary <- heart_test_predictions %>%
metrics(truth = HeartDisease, estimate = .pred_class)
heart_summary

```

```

      .metric      .estimator  .estimate
A tibble: 2 × 3   <chr>      <chr>      <dbl>
1 accuracy      binary    0.8296943
2 kap           binary    0.6569354

```

4 Expected outcomes and significance:

We expect to find that the 4 predictor variables: age, sex, chest pain type, and maximum heart rate will help us tell if a person has heart disease or not, as there's a strong correlation between each predictor variable and heart disease. The impact of these findings would be very significant. People with cardiovascular disease or who are at high cardiovascular risk (due to the presence of one or more risk factors such as hypertension, diabetes, hyperlipidemia or already established disease) need early detection and management wherein a machine learning model can be of great help. In the future, we should try to improve the accuracy of the overall algorithm by adding more predictor variables and having a large sample scale for the data. This same technique can then be applied to detect other types of diseases such as pneumonia in patients. The algorithm removes human error and there's very little chance of the algorithm misdiagnosing someone if it has a strong accuracy.

[]: